

6/14

rian, W.N. Street and W.H. Wolberg. Breast cancer diagnosis and prognosis via linear programming. Operations Research, 43(4), pages 570–577, July–August 1995. – W.H. Wolberg, W.N. Street, and O.L. Mangasarian. Machine learning techniques to diagnose breast cancer from fine-needle aspirates. Cancer Letters 77 (1994) 163–171.

```
'feature_names': array(['mean radius', 'mean texture', 'mean perimeter', 'mean area',
                        'mean smoothness', 'mean compactness', 'mean concavity',
                        'mean concave points', 'mean symmetry', 'mean fractal dimension',
                        'radius error', 'texture error', 'perimeter error', 'area error',
                        'smoothness error', 'compactness error', 'concavity error',
                        'concave points error', 'symmetry error',
                        'fractal dimension error', 'worst radius', 'worst texture',
                        'worst perimeter', 'worst area', 'worst smoothness',
                        'worst compactness', 'worst concavity', 'worst concave points',
                        'worst symmetry', 'worst fractal dimension'], dtype='<U23'),
'filename': 'breast_cancer.csv',
'data_module': 'sklearn.datasets.data'}
```

```
In [205]: train_input , test_input , train_target , test_target = train_test_split(cancer.data, cancer.target , random_state = 42)
```

executed in 15ms, finished 11:57:24 2023-10-25

```
In [206]: from sklearn.preprocessing import StandardScaler
ss = StandardScaler()
ss.fit(train_input)
```

executed in 13ms, finished 11:57:24 2023-10-25

Out[206]:

```
▼ StandardScaler
StandardScaler()
```

```
In [207]: train_scaled = ss.transform(train_input)
test_scaled = ss.transform(test_input)
```

executed in 6ms, finished 11:57:24 2023-10-25

```
In [208]: decisions = lr.decision_function(train_scaled[:, :5])
decisions
```

executed in 17ms, finished 11:57:24 2023-10-25

Out[208]:

```
array([[ 0.07042131,  5.29486531, -8.44359052, ...,  6.59081425,
         1.39492013, -2.67012852],
       [-3.8751889 , -4.37157959, 11.31746922, ...,  2.58559666,
        -9.40217833,  4.16025488],
       [-1.1130764 ,  2.08609317, -0.4867963 , ...,  0.13192612,
         1.92713095, -4.4936687 ],
       ...,
       [ 2.45467921,  5.18689785, -7.73469202, ...,  1.55121768,
         2.43239063, -5.60249733],
       [-2.31393068, -2.85766968,  7.74188303, ...,  3.46854077,
        -9.03706072,  3.65339436],
       [-4.079075 , -0.09600672,  7.38259626, ..., -3.22697538,
         1.59404047, -4.16571357]])
```

```
In [209]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression()
lr.fit(train_scaled , train_target)
```

executed in 22ms, finished 11:57:24 2023-10-25

Out[209]:

```
▼ LogisticRegression
LogisticRegression()
```

```
In [210]: lr.score(train_scaled , train_target)
```

executed in 7ms, finished 11:57:24 2023-10-25

Out[210]: 0.9859154929577465

```
In [211]: lr.score(test_scaled , test_target)
```

executed in 15ms, finished 11:57:24 2023-10-25

Out[211]: 0.9790209790209791

```
In [212]: lr.predict(test_scaled[:5])
```

executed in 14ms, finished 11:57:24 2023-10-25

Out[212]: array([1, 0, 0, 1, 1])

```
In [213]: test_scaled.shape
```

```
executed in 8ms, finished 11:57:24 2023-10-25
```

```
Out[213]: (143, 30)
```

```
In [214]: proba = lr.predict_proba(test_scaled[:5])  
np.round(proba , 3)
```

```
executed in 13ms, finished 11:57:24 2023-10-25
```

```
Out[214]: array([[0.124, 0.876],  
                [1.    , 0.    ],  
                [0.997, 0.003],  
                [0.001, 0.999],  
                [0.    , 1.    ]])
```


```
In [215]: train_score = []  
test_score = []  
for i in range(1,30):  
    lr = LogisticRegression(C = i)  
    lr.fit(train_scaled , train_target)  
    train_score.append(lr.score(train_scaled , train_target))  
    test_score.append(lr.score(test_scaled , test_target))
```

```
executed in 371ms, finished 11:57:24 2023-10-25
```

```
In [216]: plt.plot(train_score)  
plt.plot(test_score)
```

```
executed in 263ms, finished 11:57:24 2023-10-25
```

```
Out[216]: [<matplotlib.lines.Line2D at 0x2bc84bef190>]
```



In [217]: `pd.DataFrame(cancer.data , columns = cancer.feature_names)`

executed in 35ms, finished 11:57:24 2023-10-25

Out[217]:

	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	mean fractal dimension	...	worst radius	worst texture	p
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.30010	0.14710	0.2419	0.07871	...	25.380	17.33	
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.08690	0.07017	0.1812	0.05667	...	24.990	23.41	
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.19740	0.12790	0.2069	0.05999	...	23.570	25.53	
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.24140	0.10520	0.2597	0.09744	...	14.910	26.50	
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.19800	0.10430	0.1809	0.05883	...	22.540	16.67	
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
564	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.24390	0.13890	0.1726	0.05623	...	25.450	26.40	
565	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.14400	0.09791	0.1752	0.05533	...	23.690	38.25	
566	16.60	28.08	108.30	858.1	0.08455	0.10230	0.09251	0.05302	0.1590	0.05648	...	18.980	34.12	
567	20.60	29.33	140.10	1265.0	0.11780	0.27700	0.35140	0.15200	0.2397	0.07016	...	25.740	39.42	
568	7.76	24.54	47.92	181.0	0.05263	0.04362	0.00000	0.00000	0.1587	0.05884	...	9.456	30.37	

569 rows × 30 columns

In [218]: `cancer.feature_names[:10]`

executed in 19ms, finished 11:57:24 2023-10-25

Out[218]: `array(['mean radius', 'mean texture', 'mean perimeter', 'mean area',  
 'mean smoothness', 'mean compactness', 'mean concavity',  
 'mean concave points', 'mean symmetry', 'mean fractal dimension'],  
 dtype='<U23')`

In [219]: `cancer.data[:, :10]`

executed in 13ms, finished 11:57:24 2023-10-25

Out[219]: `array([[1.799e+01, 1.038e+01, 1.228e+02, ..., 1.471e-01, 2.419e-01,  
 7.871e-02],  
 [2.057e+01, 1.777e+01, 1.329e+02, ..., 7.017e-02, 1.812e-01,  
 5.667e-02],  
 [1.969e+01, 2.125e+01, 1.300e+02, ..., 1.279e-01, 2.069e-01,  
 5.999e-02],  
 ...,  
 [1.660e+01, 2.808e+01, 1.083e+02, ..., 5.302e-02, 1.590e-01,  
 5.648e-02],  
 [2.060e+01, 2.933e+01, 1.401e+02, ..., 1.520e-01, 2.397e-01,  
 7.016e-02],  
 [7.760e+00, 2.454e+01, 4.792e+01, ..., 0.000e+00, 1.587e-01,  
 5.884e-02]])`

In [220]: `train_input , test_input , train_target , test_target = train_test_split(cancer.data[:, :10], cancer.target , random_st`

executed in 10ms, finished 11:57:24 2023-10-25

In [221]: `from sklearn.preprocessing import StandardScaler  
ss = StandardScaler()  
ss.fit(train_input)`

executed in 9ms, finished 11:57:24 2023-10-25

Out[221]:

▼ StandardScaler  
StandardScaler()

In [222]: `train_scaled = ss.transform(train_input)  
test_scaled = ss.transform(test_input)`

executed in 16ms, finished 11:57:24 2023-10-25

```
In [223]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(C = 6)
lr.fit(train_scaled , train_target)
```

executed in 27ms, finished 11:57:25 2023-10-25

Out[223]:

```
▼ LogisticRegression
LogisticRegression(C=6)
```

```
In [224]: print(lr.score(train_scaled , train_target))
print(lr.score(test_scaled , test_target))
```

executed in 14ms, finished 11:57:25 2023-10-25

```
0.9436619718309859
0.9440559440559441
```

```
In [225]: result = lr.predict(test_scaled)
from sklearn.metrics import classification_report
print(classification_report(result , test_target))
```

executed in 19ms, finished 11:57:25 2023-10-25

	precision	recall	f1-score	support
0	0.94	0.91	0.93	56
1	0.94	0.97	0.95	87
accuracy			0.94	143
macro avg	0.94	0.94	0.94	143
weighted avg	0.94	0.94	0.94	143

In [ ]:

```
In [226]: from sklearn.neighbors import KNeighborsClassifier
kn = KNeighborsClassifier()
kn.fit(train_scaled , train_target)
```

executed in 18ms, finished 11:57:25 2023-10-25

Out[226]:

```
▼ KNeighborsClassifier
KNeighborsClassifier()
```

```
In [227]: len(train_scaled)
```

executed in 16ms, finished 11:57:25 2023-10-25

Out[227]: 426


```
In [228]: train_score = []
test_score = []
for i in range(1,101):
    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(train_scaled , train_target)
    train_score.append(knn.score(train_scaled , train_target))
    test_score.append(knn.score(test_scaled , test_target))
```

executed in 1.80s, finished 11:57:26 2023-10-25

```
In [229]: plt.plot(train_score)
plt.plot(test_score)
```

executed in 127ms, finished 11:57:27 2023-10-25

Out[229]: [ <matplotlib.lines.Line2D at 0x2bc830126a0>]



```
In [230]: for i in range(len(train_score)):
           if train_score[i] < test_score[i]:
               print(i)
               break
```

executed in 4ms, finished 11:57:27 2023-10-25

6

```
In [231]: from sklearn.neighbors import KNeighborsClassifier
kn = KNeighborsClassifier(n_neighbors = 5)
kn.fit(train_scaled , train_target)
```

executed in 16ms, finished 11:57:27 2023-10-25

Out[231]: 

▼ KNeighborsClassifier

KNeighborsClassifier()

```
In [232]: print(kn.score(train_scaled , train_target))
           print(kn.score(test_scaled , test_target))
```

executed in 26ms, finished 11:57:27 2023-10-25

0.9553990610328639  
0.9440559440559441

```
In [233]: from sklearn.datasets import load_iris
iris = load_iris()
```

executed in 14ms, finished 11:57:27 2023-10-25



```
In [239]: from sklearn.linear_model import LogisticRegression
lr = LogisticRegression(C = 10)
lr.fit(train_input , train_target)
print(lr.score(train_input , train_target))
print(lr.score(test_input , test_target))
```

executed in 20ms, finished 11:57:27 2023-10-25

0.9910714285714286  
1.0

```
In [240]: np.round(lr.predict_proba(train_input) , 3)
```

executed in 15ms, finished 11:57:27 2023-10-25

```
Out[240]: array([[0.997, 0.003, 0.   ],
 [0.998, 0.002, 0.   ],
 [0.   , 0.036, 0.964],
 [0.001, 0.961, 0.038],
 [0.   , 0.962, 0.038],
 [0.998, 0.002, 0.   ],
 [0.994, 0.006, 0.   ],
 [0.004, 0.995, 0.001],
 [0.   , 0.496, 0.504],
 [0.   , 0.054, 0.946],
 [0.   , 0.992, 0.007],
 [0.   , 0.001, 0.999],
 [0.001, 0.991, 0.008],
 [0.   , 0.   , 1.   ],
 [0.001, 0.908, 0.092],
 [0.996, 0.004, 0.   ],
 [0.   , 0.   , 1.   ],
 [0.001, 0.998, 0.001],
 [0.994, 0.006, 0.   ],
 [0.998, 0.002, 0.   ]]
```

```
In [241]: decisions = lr.decision_function(train_input)
decisions
```

executed in 18ms, finished 11:57:27 2023-10-25

```
Out[241]: array([[ 12.30534073,   6.49706033, -18.80240106],
 [ 13.01663648,   6.89327622, -19.9099127 ],
 [-7.89444059,   2.30269426,   5.59174633],
 [-3.57225913,   3.40023233,   0.1720268 ],
 [-5.07740319,   4.15170857,   0.92569462],
 [ 12.82031872,   6.55558949, -19.37590821],
 [ 11.95354448,   6.84701483, -18.80055931],
 [-1.09875917,   4.33362381,  -3.23486465],
 [-7.16806645,   3.57625482,   3.59181163],
 [-8.75562338,   2.94667037,   5.80895301],
 [-3.84484923,   4.36656271,  -0.52171348],
 [-11.71768411,   2.13449051,   9.58319361],
 [-2.72431581,   3.77819383,  -1.05387802],
 [-14.41227508,   3.03730925,  11.37496583],
 [-4.06001417,   3.17599955,   0.88401463],
 [ 12.06535383,   6.58233002, -18.64768385],
 [-15.0856491 ,   3.26829161,  11.81735749],
 [-2.22926363,   4.48499912,  -2.2557355 ],
 [ 11.82536693,   6.66759971, -18.49296664],
 [ 10.22252289,   6.2212557 , -18.66222289]]
```

```
In [242]: from sklearn.neighbors import KNeighborsClassifier
kn = KNeighborsClassifier(n_neighbors = 3)
kn.fit(train_input , train_target)
print(kn.score(train_input , train_target))
print(kn.score(test_input , test_target))
```

executed in 13ms, finished 11:57:27 2023-10-25

0.9464285714285714  
1.0

```
In [243]: train_score = []
test_score = []
for i in range(1,20):
    knn = KNeighborsClassifier(n_neighbors = i)
    knn.fit(train_input , train_target)
    train_score.append(knn.score(train_input , train_target))
    test_score.append(knn.score(test_input , test_target))
```


executed in 112ms, finished 11:57:27 2023-10-25



```
In [244]: plt.plot(train_score)
plt.plot(test_score)
```

executed in 132ms, finished 11:57:27 2023-10-25

Out[244]: [<matplotlib.lines.Line2D at 0x2bc83b97550>]



```
In [245]: from sklearn.neighbors import KNeighborsClassifier
kn = KNeighborsClassifier(n_neighbors = 7)
kn.fit(train_input , train_target)
print(kn.score(train_input , train_target))
print(kn.score(test_input , test_target))
```

executed in 18ms, finished 11:57:27 2023-10-25

0.9732142857142857  
1.0

```
In [246]: result = kn.predict(test_input)
from sklearn.metrics import classification_report
print(classification_report(result , test_target))
```

executed in 15ms, finished 11:57:27 2023-10-25

	precision	recall	f1-score	support
0	1.00	1.00	1.00	15
1	1.00	1.00	1.00	11
2	1.00	1.00	1.00	12
accuracy			1.00	38
macro avg	1.00	1.00	1.00	38
weighted avg	1.00	1.00	1.00	38

```
In [247]: iris.target_names[result]
```

executed in 15ms, finished 11:57:27 2023-10-25

Out[247]: array(['versicolor', 'setosa', 'virginica', 'versicolor', 'versicolor',  
'setosa', 'versicolor', 'virginica', 'versicolor', 'versicolor',  
'virginica', 'setosa', 'setosa', 'setosa', 'setosa', 'versicolor',  
'virginica', 'versicolor', 'versicolor', 'virginica', 'setosa',  
'virginica', 'setosa', 'virginica', 'virginica', 'virginica',  
'virginica', 'virginica', 'setosa', 'setosa', 'setosa', 'setosa',  
'versicolor', 'setosa', 'setosa', 'virginica', 'versicolor',  
'setosa'], dtype='<U10')

In [ ]: